

Project 2

January 26, 2024

```
[1]: # Importing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier

from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.tree import export_graphviz
from sklearn.metrics import roc_auc_score

from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler

from collections import Counter
from IPython.display import Image
import pydotplus

# Configuring the notebook
sns.set()
%matplotlib inline

# Reading the data
data = pd.read_csv('creditcard.csv')
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[2]: data.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

```
[3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    284807 non-null   float64
1   V1      284807 non-null   float64
2   V2      284807 non-null   float64
3   V3      284807 non-null   float64
4   V4      284807 non-null   float64
5   V5      284807 non-null   float64
6   V6      284807 non-null   float64
7   V7      284807 non-null   float64
8   V8      284807 non-null   float64
9   V9      284807 non-null   float64
10  V10     284807 non-null   float64
11  V11     284807 non-null   float64
```

```

12 V12      284807 non-null float64
13 V13      284807 non-null float64
14 V14      284807 non-null float64
15 V15      284807 non-null float64
16 V16      284807 non-null float64
17 V17      284807 non-null float64
18 V18      284807 non-null float64
19 V19      284807 non-null float64
20 V20      284807 non-null float64
21 V21      284807 non-null float64
22 V22      284807 non-null float64
23 V23      284807 non-null float64
24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount    284807 non-null float64
30 Class     284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```
[4]: data[['Time', 'Amount']].describe()
```

```

[4]:
count    284807.000000    284807.000000
mean      94813.859575      88.349619
std       47488.145955     250.120109
min         0.000000         0.000000
25%       54201.500000         5.600000
50%       84692.000000        22.000000
75%      139320.500000        77.165000
max      172792.000000     25691.160000

```

```
[5]: data.isnull().sum()
```

```

[5]: Time      0
V1           0
V2           0
V3           0
V4           0
V5           0
V6           0
V7           0
V8           0
V9           0
V10          0
V11          0

```

```
V12      0
V13      0
V14      0
V15      0
V16      0
V17      0
V18      0
V19      0
V20      0
V21      0
V22      0
V23      0
V24      0
V25      0
V26      0
V27      0
V28      0
Amount    0
Class     0
dtype: int64
```

```
[6]: data.fillna(data.mean(), inplace=True)
```

```
[7]: z_scores = (data - data.mean()) / data.std()
outliers = (np.abs(z_scores) > 3).any(axis=1)
print("Number of outliers:", outliers.sum())
```

Number of outliers: 37864

```
[8]: fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(15,8))
sns.distplot(data['Time'][data['Class'] == 1], bins=15, ax=ax1)
sns.distplot(data['Time'][data['Class'] == 0], bins=15, ax=ax2)

sns.distplot(data['Amount'][data['Class'] == 1], bins=5, ax=ax3)
sns.distplot(data['Amount'][data['Class'] == 0], bins=5, ax=ax4)

ax1.set_title('Fraud')
ax2.set_title('Non Fraud')
ax3.set_title('Fraud')
ax4.set_title('Non Fraud')
plt.tight_layout()
plt.show()
```

C:\Users\leend\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for

histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\leend\anaconda3\lib\site-packages\seaborn\distributions.py:2619:

FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\leend\anaconda3\lib\site-packages\seaborn\distributions.py:2619:

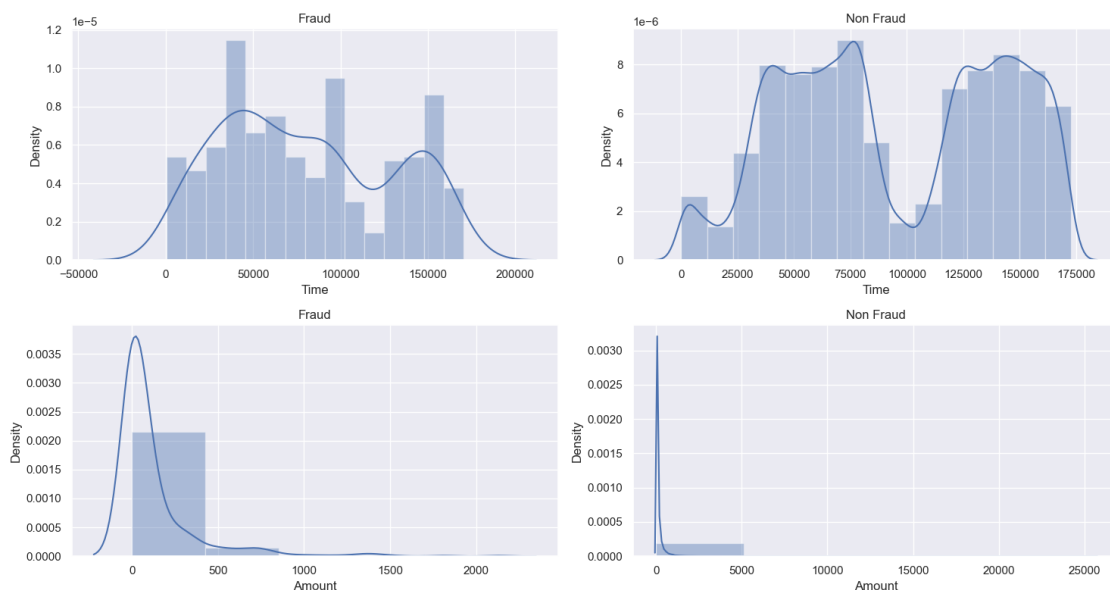
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\leend\anaconda3\lib\site-packages\seaborn\distributions.py:2619:

FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



```
[9]: fig, ax = plt.subplots(figsize=(15,8))

box_plot = sns.boxplot(data['Class'], data['Amount'], showmeans=True, ax=ax)
plt.xticks([0, 1], ['Non-Fraud', 'Fraud'])
ax.set_ylim(0, 500)
ax.set_title('Amount Distribution')
ax.set_xlabel('')
```

```

means = data.groupby(['Class'])['Amount'].mean()
vertical_offset = data['Amount'].mean() - 130 # offset from median for display

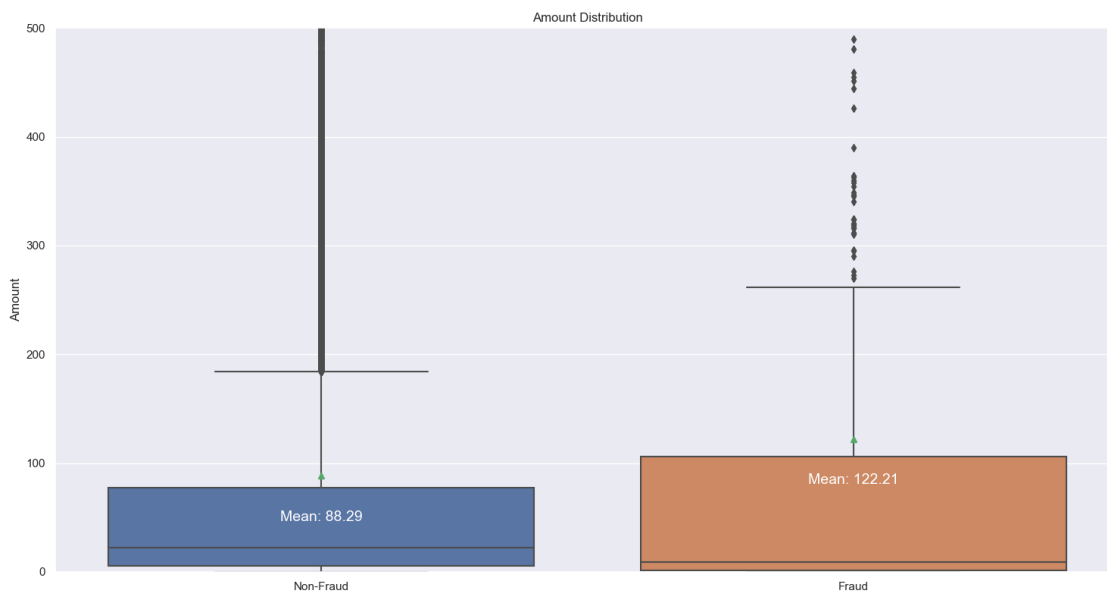
for xtick in box_plot.get_xticks():
    box_plot.text(xtick, means[xtick] + vertical_offset, f'Mean:␣
↳{round(means[xtick], 2)}',
                  horizontalalignment='center', fontsize=14, color='w')

plt.tight_layout()
plt.show()

```

C:\Users\leend\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```




```

[10]: corr = data.corr()
fig, ax = plt.subplots(figsize=(10, 8))
sns.heatmap(corr, cmap='RdBu', vmin=-1)

plt.tight_layout()
plt.show()

```



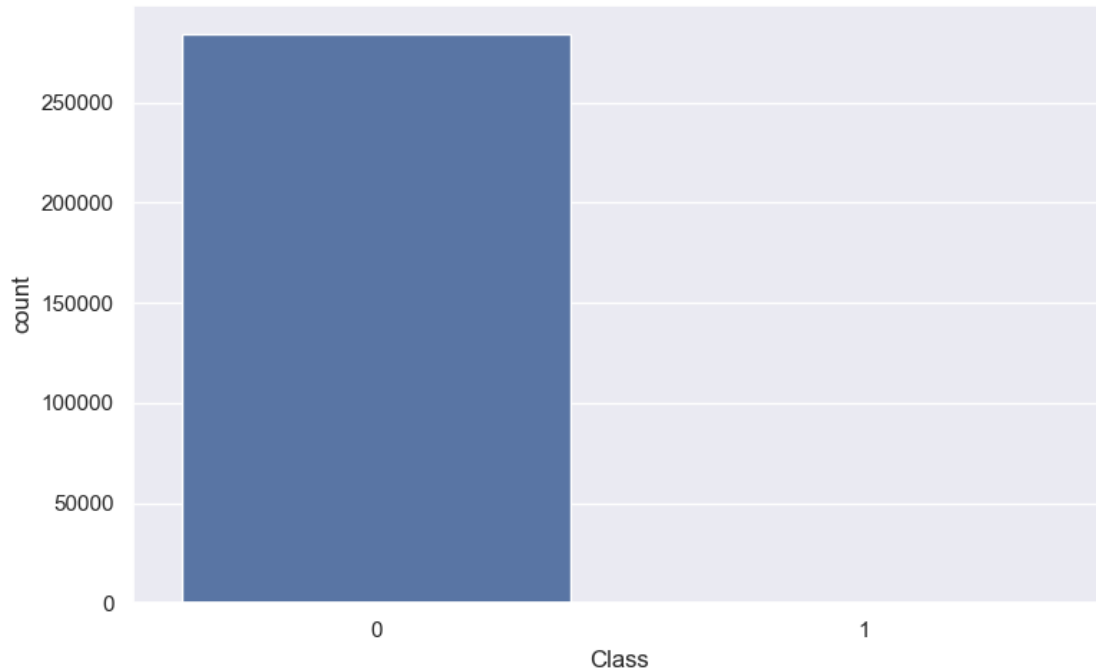
```
[11]: print(data['Class'].value_counts(normalize=True))
```

```
fig, ax = plt.subplots(figsize=(8,5))
sns.countplot(data['Class'])

plt.tight_layout()
plt.show()
```

```
0    0.998273
1    0.001727
Name: Class, dtype: float64
```

```
C:\Users\leend\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```



```
[12]: scaler = StandardScaler()

data['std_amount'] = scaler.fit_transform(data['Amount'].values.reshape(-1, 1))
data['std_time'] = scaler.fit_transform(data['Time'].values.reshape(-1, 1))

data.drop(['Amount', 'Time'], axis=1, inplace=True)
data.head()
```

```
[12]:
```

	V1	V2	V3	V4	V5	V6	V7 \
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941

	V8	V9	V10	...	V22	V23	V24	V25 \
0	0.098698	0.363787	0.090794	...	0.277838	-0.110474	0.066928	0.128539
1	0.085102	-0.255425	-0.166974	...	-0.638672	0.101288	-0.339846	0.167170
2	0.247676	-1.514654	0.207643	...	0.771679	0.909412	-0.689281	-0.327642
3	0.377436	-1.387024	-0.054952	...	0.005274	-0.190321	-1.175575	0.647376
4	-0.270533	0.817739	0.753074	...	0.798278	-0.137458	0.141267	-0.206010

	V26	V27	V28	Class	std_amount	std_time
0	-0.189115	0.133558	-0.021053	0	0.244964	-1.996583


```

1  0.125895 -0.008983  0.014724      0  -0.342475 -1.996583
2 -0.139097 -0.055353 -0.059752      0   1.160686 -1.996562
3 -0.221929  0.062723  0.061458      0   0.140534 -1.996562
4  0.502292  0.219422  0.215153      0  -0.073403 -1.996541

```

[5 rows x 31 columns]

```

[13]: X = data.drop('Class', axis=1)
      y = data['Class']

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

```

```

[14]: from imblearn.under_sampling import RandomUnderSampler

```

```

# Create a RandomUnderSampler object
rus = RandomUnderSampler()

# Use fit_resample instead of fit_sample
X_rus, y_rus = rus.fit_resample(X_train, y_train)

# Check the class distribution after undersampling
print(pd.Series(y_rus).value_counts(normalize=True))
fig, ax = plt.subplots(figsize=(8,5))
sns.countplot(y_rus)

plt.tight_layout()
plt.show()

```

```


0    0.5
1    0.5
Name: Class, dtype: float64

```

```

C:\Users\leend\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(

```



```
[15]: from imblearn.over_sampling import RandomOverSampler
```

```
# Create a RandomUnderSampler object
ros = RandomOverSampler()


# Use fit_resample instead of fit_sample
X_ros, y_ros = ros.fit_resample(X_train, y_train)

# Check the class distribution after undersampling
print(pd.Series(y_ros).value_counts(normalize=True))
fig, ax = plt.subplots(figsize=(8,5))
sns.countplot(y_ros)

plt.tight_layout()
plt.show()
```

```
0    0.5
1    0.5
Name: Class, dtype: float64
```

```
C:\Users\leend\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```



```
[16]: # logistic regression
log_reg_under = LogisticRegression()

log_reg_under.fit(X_rus, y_rus)

y_pred_log_under = log_reg_under.predict(X_test)

def report(pred):
    print(classification_report(y_test, pred))
    fig, ax = plt.subplots()
    sns.heatmap(confusion_matrix(y_test, pred, normalize='true'), annot=True,
    ↪ax=ax)
    ax.set_title('Confusion Matrix')
    ax.set_ylabel('True')
    ax.set_xlabel('Predicted')


    plt.show()

    print(f'ROC AUC Score: {round(roc_auc_score(y_test, pred), 4)}')

report(y_pred_log_under)
```

```
precision    recall  f1-score   support
```

	0	1.00	0.96	0.98	85291
	1	0.04	0.92	0.08	152
accuracy				0.96	85443
macro avg		0.52	0.94	0.53	85443
weighted avg		1.00	0.96	0.98	85443



ROC AUC Score: 0.9411

```
[17]: log_reg_over = LogisticRegression()

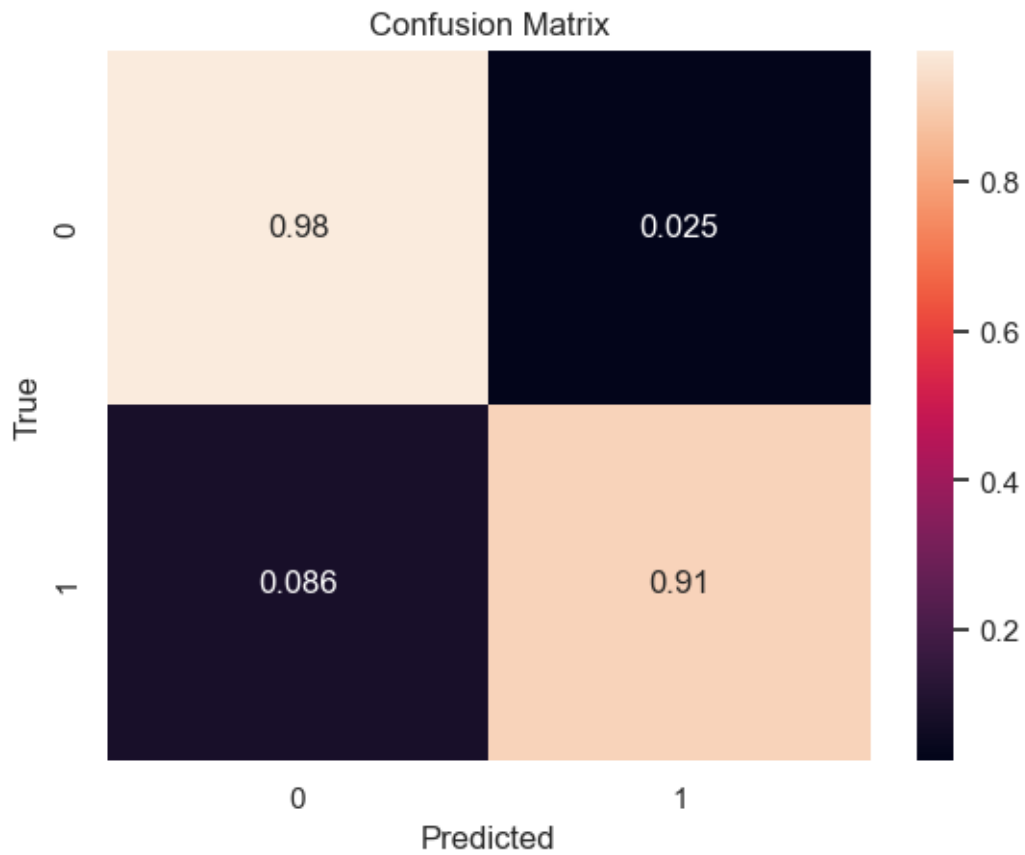
log_reg_over.fit(X_ros, y_ros)

y_pred_log_over = log_reg_over.predict(X_test)

report(y_pred_log_over)
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	1.00	0.98	0.99	85291
	1	0.06	0.91	0.12	152
accuracy				0.98	85443
macro avg		0.53	0.94	0.55	85443
weighted avg		1.00	0.98	0.99	85443



ROC AUC Score: 0.9448

```
[18]: # decision tree
n = 11
acc_tree = np.zeros((n-3))

for i in range(3, n):

    tree = DecisionTreeClassifier(criterion='entropy', max_depth=i)

    tree.fit(X_rus, y_rus)
```

```

y_pred_tree = tree.predict(X_test)

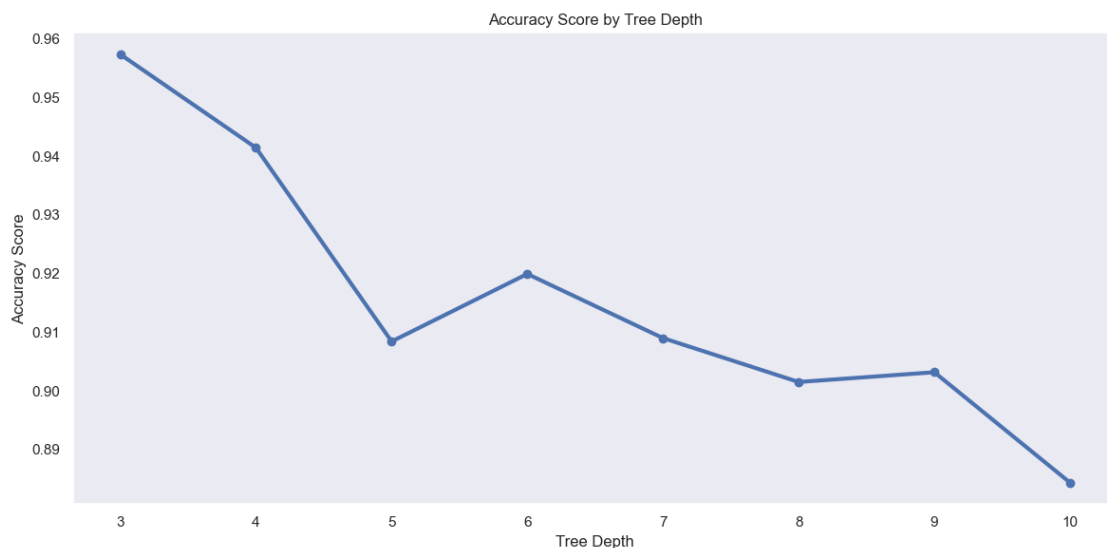
acc_tree[i-3] = accuracy_score(y_test, y_pred_tree)

fig, ax = plt.subplots(figsize=(12,6))
ax.plot(range(3, n), acc_tree, linewidth=3, marker='o')
ax.set_title('Accuracy Score by Tree Depth')
ax.set_ylabel('Accuracy Score')
ax.set_xlabel('Tree Depth')
ax.grid(False)

plt.tight_layout()
plt.show()

best_depth = acc_tree.argmax()+3
print(f'The best accuracy was {round(acc_tree.max(), 4)} with_
↳depth={best_depth}.')

```



The best accuracy was 0.9574 with depth=3.

```

[19]: tree_under = DecisionTreeClassifier(criterion='entropy', max_depth=best_depth)


tree_under.fit(X_rus, y_rus)

y_pred_tree_under = tree_under.predict(X_test)

report(y_pred_tree_under)

```

	precision	recall	f1-score	support
0	1.00	0.96	0.98	85291
1	0.04	0.89	0.07	152
accuracy			0.96	85443
macro avg	0.52	0.92	0.52	85443
weighted avg	1.00	0.96	0.98	85443




ROC AUC Score: 0.9228

```
[20]: tree_over = DecisionTreeClassifier(criterion='entropy', max_depth=best_depth)
      tree_over.fit(X_ros, y_ros)
      y_pred_tree_over = tree_over.predict(X_test)
      report(y_pred_tree_over)
```

precision	recall	f1-score	support
-----------	--------	----------	---------

0	1.00	0.93	0.97	85291
1	0.02	0.91	0.05	152
accuracy			0.93	85443
macro avg	0.51	0.92	0.51	85443
weighted avg	1.00	0.93	0.96	85443



ROC AUC Score: 0.9245

```
[21]: # K - nearest Neighbors
Ks = 11
acc_knn = np.zeros((Ks-1))

for k in range(1, Ks):

    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_rus, y_rus)
```



```

y_pred_knn = knn.predict(X_test)

acc_knn[k-1] = accuracy_score(y_test, y_pred_knn)

fig, ax = plt.subplots(figsize=(12,6))
ax.plot(range(1, Ks), acc_knn, linewidth=3, marker='o')
ax.set_title('Accuracy Score by Number of Neighbors')
ax.set_ylabel('Accuracy Score')
ax.set_xlabel('Number of Neighbors')
ax.grid(False)

plt.tight_layout()
plt.show()

best_k = acc_knn.argmax()+1
print(f'The best accuracy was {round(acc_knn.max(), 4)} with k={best_k}.')

```

C:\Users\leend\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\leend\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\leend\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\leend\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```

mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\leend\anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other
reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode`
typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will
change: the default value of `keepdims` will become False, the `axis` over which
the statistic is taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\leend\anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other
reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode`
typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will
change: the default value of `keepdims` will become False, the `axis` over which
the statistic is taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\leend\anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other
reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode`
typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will
change: the default value of `keepdims` will become False, the `axis` over which
the statistic is taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\leend\anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other
reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode`
typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will
change: the default value of `keepdims` will become False, the `axis` over which
the statistic is taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\leend\anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other
reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode`
typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will
change: the default value of `keepdims` will become False, the `axis` over which
the statistic is taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\leend\anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other
reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode`
typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will
change: the default value of `keepdims` will become False, the `axis` over which
the statistic is taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.

```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```



The best accuracy was 0.9882 with k=10.

```
[22]: knn_under = KNeighborsClassifier(n_neighbors=best_k)

knn_under.fit(X_rus, y_rus)


y_pred_knn_under = knn_under.predict(X_test)

report(y_pred_knn_under)
```

C:\Users\leend\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

	precision	recall	f1-score	support
0	1.00	0.99	0.99	85291
1	0.12	0.88	0.21	152
accuracy			0.99	85443
macro avg	0.56	0.93	0.60	85443
weighted avg	1.00	0.99	0.99	85443



ROC AUC Score: 0.935

```
[24]: knn_over = KNeighborsClassifier(n_neighbors=best_k)

knn_over.fit(X_ros, y_ros)

y_pred_knn_over = knn_over.predict(X_test)


report(y_pred_knn_over)
```

C:\Users\leend\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
precision    recall  f1-score   support
```

	0	1.00	1.00	1.00	85291
	1	0.55	0.84	0.66	152
accuracy				1.00	85443
macro avg		0.77	0.92	0.83	85443
weighted avg		1.00	1.00	1.00	85443



ROC AUC Score: 0.9204

```
[25]: #XGBoost


xgb = XGBClassifier()

xgb.fit(X_train, y_train)

y_pred_xgb = xgb.predict(X_test)

report(y_pred_xgb)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85291
1	0.92	0.80	0.86	152
accuracy			1.00	85443
macro avg	0.96	0.90	0.93	85443
weighted avg	1.00	1.00	1.00	85443



ROC AUC Score: 0.9013

```
[26]: # Calculating the ratio
counter = Counter(y_train)
estimate = counter[0] / counter[1]


# Implementing the model
xgb = XGBClassifier(scale_pos_weight=estimate)

xgb.fit(X_train, y_train)
```

```
y_pred_xgb_scaled = xgb.predict(X_test)

report(y_pred_xgb_scaled)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85291
1	0.89	0.81	0.85	152
accuracy			1.00	85443
macro avg	0.95	0.90	0.92	85443
weighted avg	1.00	1.00	1.00	85443



ROC AUC Score: 0.9045

```
[27]: #Comparing the models
summary = pd.DataFrame(data={
    'labels': ['Accuracy', 'Precision', 'Recall', 'F1_score', 'roc_auc'],
```

```

'log_reg_under': [accuracy_score(y_test, y_pred_log_under),
↳precision_score(y_test, y_pred_log_under), recall_score(y_test,
↳y_pred_log_under), f1_score(y_test, y_pred_log_under), roc_auc_score(y_test,
↳y_pred_log_under)],
'log_reg_over': [accuracy_score(y_test, y_pred_log_over),
↳precision_score(y_test, y_pred_log_over), recall_score(y_test,
↳y_pred_log_over), f1_score(y_test, y_pred_log_over), roc_auc_score(y_test,
↳y_pred_log_over)],
'decision_trees_under': [accuracy_score(y_test, y_pred_tree_under),
↳precision_score(y_test, y_pred_tree_under), recall_score(y_test,
↳y_pred_tree_under), f1_score(y_test, y_pred_tree_under),
↳roc_auc_score(y_test, y_pred_tree_under)],
'decision_trees_over': [accuracy_score(y_test, y_pred_tree_over),
↳precision_score(y_test, y_pred_tree_over), recall_score(y_test,
↳y_pred_tree_over), f1_score(y_test, y_pred_tree_over), roc_auc_score(y_test,
↳y_pred_tree_over)],
'knn_under': [accuracy_score(y_test, y_pred_knn_under), precision_score(y_test,
↳y_pred_knn_under), recall_score(y_test, y_pred_knn_under), f1_score(y_test,
↳y_pred_knn_under), roc_auc_score(y_test, y_pred_knn_under)],
'knn_over': [accuracy_score(y_test, y_pred_knn_over), precision_score(y_test,
↳y_pred_knn_over), recall_score(y_test, y_pred_knn_over), f1_score(y_test,
↳y_pred_knn_over), roc_auc_score(y_test, y_pred_knn_over)],
'XGBoost': [accuracy_score(y_test, y_pred_xgb), precision_score(y_test,
↳y_pred_xgb), recall_score(y_test, y_pred_xgb), f1_score(y_test, y_pred_xgb),
↳roc_auc_score(y_test, y_pred_xgb)],
'XGBoost_scaled': [accuracy_score(y_test, y_pred_xgb_scaled),
↳precision_score(y_test, y_pred_xgb_scaled), recall_score(y_test,
↳y_pred_xgb_scaled), f1_score(y_test, y_pred_xgb_scaled),
↳roc_auc_score(y_test, y_pred_xgb_scaled)]
}).set_index('labels')
summary.index.name = None

summary

```

```

[27]:

```

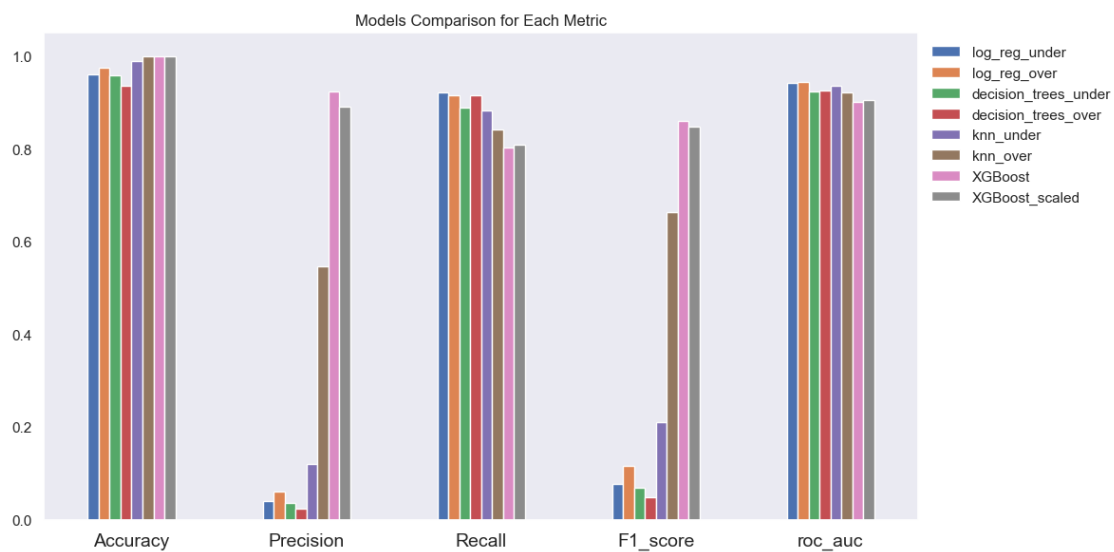
	log_reg_under	log_reg_over	decision_trees_under	\
Accuracy	0.961167	0.975013	0.957352	
Precision	0.040627	0.061477	0.035885	
Recall	0.921053	0.914474	0.888158	
F1_score	0.077821	0.115209	0.068983	
roc_auc	0.941146	0.944797	0.922816	

	decision_trees_over	knn_under	knn_over	XGBoost	XGBoost_scaled
Accuracy	0.934588	0.988179	0.998479	0.999532	0.999485
Precision	0.024322	0.119005	0.547009	0.924242	0.891304
Recall	0.914474	0.881579	0.842105	0.802632	0.809211
F1_score	0.047384	0.209703	0.663212	0.859155	0.848276

roc_auc 0.924549 0.934974 0.920431 0.901257 0.904517

```
[28]: fig, ax = plt.subplots(figsize=(12, 6))
summary.plot.bar(ax=ax)
ax.legend(bbox_to_anchor=(1, 1), frameon=False)
ax.grid(False)
ax.set_title('Models Comparison for Each Metric')

plt.xticks(rotation=0, fontsize=14)
plt.tight_layout()
plt.show()
```



```
[ ]:
```